

Titanic - Advanced Feature Engineering Tutorial.

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
sns.set(style='darkgrid')

```

} For EDA

For Model.

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import LabelEncoder, OneHotEncoder, StandardScaler
from sklearn.metrics import roc_curve, auc
from sklearn.model_selection import StratifiedKFold

```

```

import string
import warnings
warnings.filterwarnings('ignore')
SEED = 42

```

} Default.

0. Loading Data

```
def concat_df(train_data, test_data):
```

```

    # Index 32
    # Index drop
    return pd.concat([train_data, test_data], sort=True).reset_index(drop=True)

```

```
def divide_df(all_data)
```

```

    # change
    return all_data.iloc[:, :890], all_data.iloc[891:, :].drop(['Survived'], axis=1)

```

```
df_train = pd.read_csv('~')
```

```
df_test = ''
```

```
df_all = concat_df(df_train, df_test)
```

```
df_train.name = 'Training Set'
```

```
df_test.name = 'Test Set'
```

```
df_all.name = 'All Set'
```

```
dfs = [df_train, df_test]
```

각 Set의 Shape, Columns 확인.

1. EDA.

1.1 Overview the dataset

1.2 Missing Values.

```
def display_missing(df):
```

```
    for col in df.columns.tolist():
```

```
        print(f'{col} columns missing values : {df[col].isnull().sum()}')
```

```
for df in dfs:
```

```
    print(df.name)
```

```
    display_missing(df)
```


1.2.1 Age.

- How to fill NA in 'Age'? By using high correlation with other features.

```
df_all_corr = df_all.corr().abs().unstack().sort_values(kind='quicksort',
                                                    ascending=False).reset_index()
```

df_all_corr 의 type : Series → DataFrame (by 'reset_index' method)
(이름 변경) , inplace=True.

```
df_all_corr.rename(columns={'level_0': 'Feature_1', ... 3})
```

```
df_all_corr[df_all_corr['Feature_1'] == 'Age']
```

'Age'와 'Pclass'의 상관 관계가 가장 높음. 하지만 이것만으로 feature engineering
하기엔 너무 상관함. 정확도를 높이기 위해 'Sex' feature도 고려하여 missing value 처리.

groupby
age-by-pclass-sex = df_all[['Pclass', 'Sex']].median()['Age']

```
for pclass in df_all['Pclass'].unique().tolist():
```

```
    for sex in df_all['Sex'].unique().tolist():
```

```
        print(f"median age of {pclass} {sex} : ", age-by-pclass-sex[pclass][sex])
```

```
df_all['Age'] = df_all.groupby(['Pclass', 'Sex'])['Age'].apply(lambda x: x.fillna(x.median()))
```

'lambda'는 DataFrame 객체의 Index 값을 받음. 위의 경우인 (1, 'female')과
같은 이름 인덱스로 x로 받음.

1.2.2 Embarked

- Need to fill two missing values. (Using outside information)

```
df_all['Embarked'] = df_all['Embarked'].fillna('S')
```

1.2.3 Fare

- Need to fill one missing value.

- The passenger is male, third-class, and no family.

```
med_fare = df_all.groupby(['Pclass', 'Parch', 'Sibsp'])['Fare'].  
            median()[3][0][0]
```

```
df_all['Fare'] = df_all['Fare'].fillna(med_fare)
```

1.2.4 Cabin

```
df_all['Deck'] = df_all['Cabin'].apply(lambda s: s[0] if pd.notnull(s)  
                                         else 'M')
```

→ 'Cabin'의 첫 문자가 Deck을 의미함. Deck 별 Pclass 상자를 확인.

```
df_all_decks = df_all.groupby(['Pclass', 'Deck']).count().  
                drop(columns=[' ~ '])
```

```
def get_pclass_dist(df):
```

```
    deck_counts = {'A': {3, ... 3}
```

```
    decks = df.columns.levels[0]
```


for deck in decks:

```
try:
    for pclass in range(1, 4):
        count = df[deck][pclass][0]
        deck_counts[deck][pclass] = count
```

except KeyError:

```
    deck_counts[deck][pclass] = 0.
```

```
df_deck = pd.DataFrame(deck_counts)
```

```
df deck_percentage = {}
```

```
for col in df_deck.columns:
```

```
    deck_percentage[col] = [(count / df_deck[col].sum()) * 100
```

```
        for count in df_deck[col]]
```

```
return deck_counts, deck_percentage.
```

```
def display_pclass_dist(percentages):
```

```
    df_percentages = pd.DataFrame(percentages).transpose()
```

```
    deck_names = ('A', ..., 'T')
```

```
    bar_count = np.arange(1/en(deck_names)) } # X ticks labeled with 1 to 16.
```

```
    bar_width = 0.85
```

```
    pclass1 = df_percentages[0]
```

```
    " 2 " [1]
```

```
    " 3 " [2]
```

```
plt.figure(figsize=(20, 10))
```

```
plt.bar ( bar_count, pclass1, color = ,
          width = bar_width, edgecolor = ,
          label = 'pclass1' )
```

```
plt.bar ( bar_count, pclass2, bottom = pclass1,
          " " )
```

```
plt.bar ( bar_count, pclass3, bottom = pclass1 + pclass2,
          " " )
```

```
plt.xlabel ( 'Deck' )
```

```
plt.ylabel ( 'Pclass percentage' )
```

```
plt.xticks ( bar_count, deck_names )
```

```
plt.legend ( loc = 'Upper left', bbox_to_anchor = (1, 1), prop = {'size': 15} )
```

```
plt.title ( ' ' )
```

```
plt.show ( )
```

```
all_deck_count, all_deck_per = get_pclass_dist (df_all_decks)
display_pclass_dist (all_deck_per)
```

change 'T' deck to 'A' deck.

```
idx = df_all [df_all ['Deck'] == 'T'].index
```

```
df_all.loc [idx, 'Deck'] = 'A'
```

* Survived & Deck & 위의 과정의 동일성

1.3 Target Distribution.

- Target value 분포 (수치, 비율)

```
Survived_cnt = df_train['Survived'].value_counts()[1]
```

```
not_survived_cnt = df_train['Survived'].value_counts()[0]
```

```
Survived_ratio = (Survived_cnt / df_train.shape[0]) * 100
```

```
not_survived_ratio = (Not_ " " ) * 100
```

```
plt.figure(figsize=(10,8))
```

```
sns.countplot(df_train['Survived'])
```

```
plt.xticks((0,1), [f'Not Survived ( {Survived_ratio:.2f} %)', f'Survived_ratio ( {not_ " " } % )'])
```

```
plt.title ( " " )
```

```
plt.show ( " " )
```

1.4 Correlations.

```
df_train_corr = df_train.drop(['PassengerId'], axis=1).corr(), abs(),  
unstacked(), sort_values(kind='quicksort', ascending=False),  
reset_index()
```

```
df_train_corr.rename(columns = {'eael_o': 'feature 1', ... 3, inplace=True)
```

```
df_train_corr.drop(df_train_corr.iloc[1:2], index, inplace=True)
```

```
" (df_train_corr[df_train_corr['Correlation Coefficient'] == 1].index,  
inplace=True)
```

```
df_test_corr
```

7

df-test & } corr = df-train-corr ['Correlation Coefficient'] > 0.1
 sol. { df-train-corr [corr]

Heatmap

f, ax = plt.figure (rows = 2, figsize = (20, 20))

sns.heatmap (df-train.drop (['PassengerId'], axis=1).corr (),

ax = ax[0], annot = True, square = True,

annot_kws = {'size': 14})

ticks label } sns.heatmap (df-test.drop (" ")

plt.show ()

1.5 Target Distribution in Features

1.5.1 Continuous Features - 'Fare' & 'Age'

Cont-features = ['Fare', 'Age']

surv = df-train ['Survived'] == 1

f, ax = plt.subplots (rows=2, ncols=2, figsize=(20,20))

plt.subplots_adjust (right = 1.5)

for i, feature in enumerate (Cont-features):

sns.histplot (df-train[~surv][feature], ax=ax[0][i],

hist = True, label = 'Not Survived', color = ' ')

sns.histplot (df-train[surv][feature], ax=ax[0][i], " ")


```
sns.histplot(df_train[feature], hist=False, ax=ax[1][i],
              label='Training Set', color=' ')
```

```
sns.histplot(df_test[feature], label='Test Set', '')
```

329 { ticks
legend
tick_params
set_title

```
plt.show()
```

1.5.2 Categorical Features

• 'Pclass', 'Sex' features

homogeneity distributions

Why? Hasty
generalization
to say

```
Cat_features = ['Embarked', 'Parch', 'Pclass', 'Sex', 'Sibsp', 'Deck']
```

```
f, ax = plt.subplots(nrows=2, ncols=3, figsize=(20, 20))
```

```
plt.subplots_adjust(right=1.5, top=1.25)
```

```
for i, feature in enumerate(Cat_features, 1):
```

```
    plt.subplot(2, 3, i)
```

```
    sns.countplot(x=feature, data=df_train, hue='Survived')
```

330 { label
tick_params
legend
title

```
plt.show()
```

2. Feature Engineering.

2.1 Bining continuous features.

2.1.1 Fare.

- This feature is positively skewed, and survival rate is extremely high on the right end.

Quantile

`df_all['Fare'] = pd.qcut(df_all['Fare'], (13))`

Why 13? empirical data?

Visualization.

`f, ax = plt.subplots(figsize=(22, 9))`

`sns.countplot(x='Fare', hue='Survived', data=df_all)`

320 {
xlabel
ylabel
tick-params
title
legend

`plt.show()`

2.1.2 Age.

- This feature shows normal distribution.

Quantile

`df_all['Age'] = pd.qcut(df_all['Age'], 10)`

X

2.2 Frequency Encoding

• 'Family-size' feature which is derived feature from Parch, Sibsp is also able to predict the survival rate.

• Family-size 1 : Alone

2,3,4 : Small

5,6 : Medium

7~ : Large

```
df_all['Family-Size'] = df_all['Parch'] + df_all['Sibsp'] + 1
```

```
f, ax = plt.subplots(nrows=2, ncols=2, figsize=(20,20))
```

```
plt.subplots_adjust(right=1.5)
```

```
sns.barplot(x=df_all['Family-Size'].value_counts().index,
```

```
y=df_all['Survived'].values,
```

```
ax=ax[0][0])
```

```
sns.countplot(x='Family-Size', hue='Survived', data=df_all, ax=ax[0][1])
```

↳ df_test set of 'Survived' x

```
family_map = {1: 'Alone', ~ 3}
```

```
df_all['Family-Size-Grouped'] = df_all['Family-Size'].map(family_map)
```

```
plt.show()
```