Titanic - Advanced Feature Engineering Tutorial

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
sns.set(style='darkgrid')
```
} For EDA

For Model
```
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import LabelEncoder, OneHotEncoder, StandardScaler
from sklearn.metrics import roc_curve, auc
from sklearn.model_selection import StratifiedKFold
```

```
import String
import warnings
warnings.filterwarnings('ignore')
SEED = 42
```
} Default

0. Loading Data

```
def concat_df(train_data, test_data):
    return pd.concat([train_data, test_data], sort=True).reset_index(drop=True)
```
(sort=True: Index 정렬, drop=True: Index drop)

```
def divide_df(all_data):
    return all_data.loc[:890], all_data.loc[891:].drop(['Survived'], axis=1)
```
(:890 change)

```
df_train = pd.read_csv('~')
df_test = "
df_all = concat_df(df_train, df_test)

df_train.name = 'Training Set'
df_test.name = 'Test Set'
df_all.name = 'All set'

dfs = [df_train, df_test]
# 각 Set의 Shape, Columns 확인.

1. EDA.
 1.1 Overview the dataset
 1.2 Missing Values.

def display_missing(df):
    for col in df.columns.tolist():
        print(f'{col} colums missing values : {df[col].isnull().sum()}')

for df in dfs:
    print(df.name)
    display_missing(df)
```

NO.

1.2.1 Age.

• How to fill NA in 'Age'? By using high correlation with other features.

```
df_all_corr = df_all.corr().abs().unstack().sort_values(kind='quicksort',
                ascending=False).reset_index()
# df_all_corr 의 type : Series → DataFrame (by 'reset_index' method)
                (이름 안줌)                          , inplace=True.
df_all_corr.rename(columns={'level_0': 'Feature_1', ... 3})
df_all_corr[df_all_corr['Feature_1'] == 'Age']
# 'Age'와 'Pclass'의 상관 관계 가장 높음. 하지만 이것만으로 feature engineering
하기엔 너무 성급함. 정확성을 높이기 위해 'Sex' feature도 고려하여 missing value 채울.
                    groupby
age_by_pclass_sex = df_all[['Pclass', 'Sex']].median()['Age']

for pclass in df_all['Pclass'].unique().tolist():
    for sex in df_all['Sex'].unique().tolist():
        print(f'Median age of {pclass} {sex} : {age_by_pclass_sex[pclass][sex]}')

df_all['Age'] = df_all.groupby(['Pclass', 'Sex'])['Age'].apply(lambda x: x.fillna(x.median()))
# 'lambda'는 DataFrame 에서의 Index 값을 받음. 위의 경우인 (1, 'female')과
같은 이름 안붙은 X로 받음.
```

NO.

1.2.2 Embarked

- Need to fill two missing values. (Using outside information)

```
df_all['Embarked'] = df_all['Embarked'].fillna('S')
```

1.2.3 Fare

- Need to fill one missing value.
  - The passenger is male, third-class, and no family.

```
med_fare = df_all.groupby(['Pclass', 'Parch', 'Sibsp'])['Fare'].
           median()[3][0][0]
df_all['Fare'] = df_all['Fare'].fillna(med_fare)
```

1.2.4 Cabin

```
df_all['Deck'] = df_all['Cabin'].apply(lambda s : s[0] if pd.notnull(s)
                                       else 'M')
```

→ 'Cabin'의 첫 문자가 Deck을 의미함. Deck 별 Pclass 정보를 확인.

```
df_all_decks = df_all.groupby(['Pclss', 'Deck']).count().
               drop(columns=['~'])
```

```
def get_pclss_dist(df):

    deck_counts = {'A': {3, ... 3
    decks = df.columns.levels[0]
```

```
for deck in decks:
    try:
        for pclass in range(1, 4):
            count = df[deck][pclass][0]
            deck_counts[deck][pclass] = count
    except KeyError:
        deck_counts[deck][pclass] = 0.

df_deck = pd.DataFrame(deck_counts)
deck_percentage = { }

for col in df_deck.columns:
    deck_percentage[col] = [(count / df_deck[col].sum()) * 100
                            for count in df_deck[col]]

return deck_counts, deck_percentage

def display_pclass_dist(percentages):
    df_percentages = pd.DataFrame(percentages).transpose()
    deck_names = ('A', ..., 'T')
    bar_count = np.arange(len(deck_names))   # Xticks label 위치 지정.
    bar_width = 0.85

    pclass1 = df_percentages[0]
    " 2       " [1]
    " 3       " [2]

    plt.figure(figsize=(20,10))
```

NO.

```
plt.bar(bar_count, pclass1, color=    ,
        width=bar_width, edgecolor=    ,
        label='Pclass1')
plt.bar(bar_count, pclass2, bottom=pclass1,
        "                   )
plt.bar(bar_count, pclass3, bottom=pclass1+pclass2,
        "                   )

plt.xlabel('Deck')
plt.ylabel('Pclass percentage')
plt.xticks(bar_count, deck_names)
plt.legend(loc='Upper left', bbox_to_anchor=(1,1), prop={'size':15})
plt.title('    ')
plt.show()

all_deck_count, all_deck_per = get_pclass_dist(df_all_decks)
display_pclass_dist(all_deck_per)

# change 'T' deck to 'A' deck
idx = df_all[df_all['Deck']=='T'].index
df_all.loc[idx,'Deck'] = 'A'
```

* Survived & Deck 간의 관계 확인

1.3 Target Distribution

- Target value 시각화 (수치, 비율)

```
Survived_cnt = df_train['Survived'].value_counts()[1]
not_survived_cnt = df_train['Survived'].value_counts()[0]
Survived_ratio = (Survived_cnt / df_train.shape[0]) * 100
not_Survived_ratio = (Not_ "                    ) * 100

plt.figure(figsize=(10,8))
sns.countplot(df_train['Survived'])

plt.xticks((0,1), [f'Not Survived ({not_Survived_ratio:.2f}%)', f'Survived_ratio(  )'])
plt.title(          )
plt.show()
```

1.4 Correlations.

```
df_train_corr = df_train.drop(['PassengerId'], axis=1).corr().abs().
                unstacked().sort_values(kind='quicksort', ascending=False).
                reset_index()
df_train_corr.rename(columns={'level_0': 'feature 1', ... }, inplace=True)
df_train_corr.drop(df_train_corr.iloc[1::2].index, inplace=True)
"                 (df_train_corr[df_train_corr['Correlation Coefficient']==1].index,
                  inplace=True)

df_test_corr
```

7

NO.

df_test도 동일.

```
corr = df_train_corr['Correlation Coefficient'] > 0.1
df_train_corr[corr]
```

Heatmap 시각화.

```
f, ax = plt.figure(nrows=2, figsize=(20,20))

sns.heatmap(df_train.drop(['PassengerId'], axis=1).corr(),
            ax=ax[0], annot=True, square=True,
            annot_kws={'size':14})
sns.heatmap(df_test.drop(  "  )
```

ticks label 제외

```
plt.show()
```

1.5 Target Distribution in Features.

1.5.1 Continuous Features - 'Fare' & 'Age'

```
Cont_features = ['Fare', 'Age']
surv = df_train['Survived'] == 1

f, ax = plt.subplots(nrows=2, ncols=2, figsize=(20,20))
plt.subplots_adjust(right=1.5)

for i, feature in enumerate(Cont_features):

    sns.histplot(df_train[~surv][feature], ax=ax[0][i],
                 kist=True, label='Not Survived', color=' ')
    sns.histplot(df_train[surv][feature], ax=ax[0][i], "  )
```

DESIGN K

```
sns.histplot(df_train[feature], hist=False, ax=ax[i][i],
             label='Training Set', color='')
sns.histplot(df_test[feature], label='Test Set', '')
```

꾸밈 { ticks, legend, tick_params, set_title }

```
plt.show()
```

1.5.2 Categorical features

- 'Pclass', 'Sex' features — homogeneity distributions — Why? { Hasty generalization to say }

```
Cat_features = ['Embarked', 'Parch', 'Pclass', 'Sex', 'Sibsp', 'Deck']

f, ax = plt.subplots(nrows=2, ncols=3, figsize=(20, 20))
plt.subplots_adjust(right=1.5, top=1.25)

for i, feature in enumerate(Cat_features, 1):

    plt.subplot(2, 3, i)
    sns.countplot(x=feature, data=df_train, hue='Survived')
```

꾸밈 { label, tick_params, legend, title }

```
plt.show()
```

NO.

2. Feature Engineering.

2.1 Bining continuous features.

2.1.1 Fare.

- This feature is positively skewed, and survival rate is extremely high on the right end.

```
# Quantile
df_all['Fare'] = pd.qcut(df_all['Fare'], 13)
```

- Why 13? empirical data?

```
# Visualization
f, ax = plt.subplots(figsize=(22,9))
sns.countplot(x='Fare', hue='Survived', data=df_all)
```

328
- xlabel
- ylabel
- tick_params
- title
- legend

```
plt.show()
```

2.1.2 Age.

- This feature shows normal distribution.

```
# Quantile
df_all['Age'] = pd.qcut(df_all['Age'], 10)
```

NO.

2.2 Frequency Encoding

- 'Family_size' feature which is derived feature from Parch, Sibsp is also able to predict the survival rate.
- Family_size 1 : Alone
  2,3,4 : Small
  5,6 : Medium
  7~ : Large.

```
df_all['Family_Size'] = df_all['Parch'] + df_all['Sibsp'] + 1

f, ax = plt.subplots(nrows=2, ncols=2, figsize=(20,20))
plt.subplots_adjust(right=1.5)

sns.barplot(x=df_all['Family_Size'].value_counts.index,
            y=                                .values,
            ax=ax[0][0])
sns.countplot(x='Family_Size', hue='Survived', data=df_all, ax=ax[0][1])
                                  └ df_test 넣으면 'Survived' X

family_map = {1:'Alone', ~ 3
df_all['Family_Size_Grouped'] = df_all['Family_Size'].map(family_map)

plt.show()
```

2.3 Title & Is married

- 'Title' is created by extracting the prefix before 'Name'
- 'Is married' is binary feature based on 'Mrs' title.

```
df_all['Title'] = df_all['Name'].str.split(', ', expand=True)[1]
                  .str.split('.', expand=True)[0]
df_all['Is_Married'] = 0
df_all['Is_Married'].loc[df_all['Title'] == 'Mrs'] = 1

# Visualization (Before & After Grouping the title)
fig, ax = plt.subplots(nrows=2, figsize=(20,20))
sns.barplot(x=df_all['Title'].value_counts().index,
            y=                    .values, ax=ax[0])

df_all['Title'] = df_all['Title'].replace([ ~ ], 'Miss/Mrs/Ms')
                  (['    '],          )
//

plt.show()
```

2.4 Target Encoding

- Creating 'Family' feature to group passengers in the same family by using 'Name' feature.

```
def extract_surname(data):

    families = []

    for i in range(len(data)):
        name = data.iloc[i]

        if '(' in name:
            name_no_bracket = name.split('(')[0]
        else:
            name_no_bracket = name

        family = name_no_bracket.split(',')[0]

        for c in string.punctuation:
            family = family.replace(c, '').strip()

        families.append(family)

    return families

df_all['Family'] = extract_surname(df_all['Name'])
df_train, df_test = divide_df(df_all)
```

Creating 'Survival rate' & 'Survival rate_NA' features.

```
non_unique_families = [x for x in df_train['Family'].unique()\
                       if x in df_test['Family'].unique()]
non_unique_tickets = [                "                ['Ticket']                "
                                      "                ['Ticket']                "  ]
```

median ?

```
df_family_survival_rate = df_train.groupby('Family')['Survived', 'Family',
                                                     'Family_Size'].median()
df_ticket_survival_rate =                "
```

```
family_rates = {}
ticket_rates = {}
```

```
for i in range( len(df_family_survival_rate) ):
    if ① df_family_survival_rate.index[i] in non_unique_families and
       ② df_family_survival_rate.iloc[i, 1] > 1 :
        family_rates[ ① ] = ② .iloc[i, 0]
```

→ Ticket too

```
mean_survival_rate = np.mean(df_train['Survived'])

train_family_survival_rate = []
"                  _NA = []
test_  "             = []
"                  _NA = []

for i in range( len(df_train) ):

    if df_train['Family'][i] in family_rates:
        train_family_survival_rate.append(family_rates[df_train['Family'][i]])
        train_family_survival_rate_NA.append(1)
    else:
        train_family_survival_rate.append(mean_survival_rate)
        "          _NA.append(0)
```

For test set, too

7

```
df_train['Family_survival_rate'] = train_family_survival_rate
"                 _NA']  =  "      _NA
```

For test set, too

7

For 'Ticket' feature, too

7

```
for df in [df_train, df_test]:
    df['Survival_Rate'] = ( df['Ticket_Survival_Rate'] + df['Family_Survival_Rate'])/2
    df[ "     _NA'] =
```

2.5 Feature Transformation.

2.5.1 Label encoding Non-Numerical features to Numeric

- Object type : 'Embarked', 'Sex', 'Deck', 'Title', 'Family_Size_grouped'
- Category type : 'Fare', 'Age'

⇒ LabelEncoder.

```
non_numeric_features = [                ]
```

'df_all'로 transforming 하는 것이 값을 관리 측면에서 좋은점은?
(1) fit_transform

```
for df in dfs:
    for feature in non_numeric_features:
        df[feature] = LabelEncoder().fit_transform(df[feature])
```

2.5.2 One-Hot Encoding the categorical features

```
Cat_features = ['Pclass', 'Deck', 'Sex', 'Family_Size_grouped',
                'Embarked', 'Title']
encoded_features = []
```

```
for df in dfs:
    for feature in cat_features:
        encoded_feat = OneHotEncoder().fit_transform(
            df[feature].values.reshape(-1,1)).toarray()
        n = df[feature].nunique()
        cols = [f'{feature}_{i}' for i in range(1, n+1)]
        encoded_df = pd.DataFrame(encoded_feat, columns=cols)
        encoded_df.index = df.index
        encoded_features.append(encoded_df)

df_train = pd.concat([df_train, *encoded_features[:6]], axis=1)
df_test  =  "                  [6:]
```

Series → Array

2.6 Conclusion

3. Model

```
X_train = ①StandardScaler().fit_transform(df_train.drop(columns=cols))
y_train = df_train['Survived'].values
X_test  = ① "
```

drop

3.1 Random Forest

- Compare Sigle model with leaderboard model, and try to experiment hyperparameter tuning

```
Single_best_model = RandomForestClassifier( criterion='gini',
                                             n_estimators=1100, # 생성될 tree 개수
                                             max_depth=5,
                                             min_samples_split=4,
                                             min_samples_leaf=5,
                                             max_features='auto',
                                             oob_score=True,
                                             random_state=SEED,
                                             n_jobs=-1
                                             # Printing 'words' information - verbose=1 )

N=5    # KFold num.
oob=0
probs = pd.DataFrame(np.zeros( (len(X_test), N*2) ),
                     columns=['Fold{}_Prob_{}'.format(i,j)
                              for i in range(1,N+1)
                              for j in range(2) ])
importances = pd.DataFrame( np.zeros( (X_train.shape[1], N) ),
                            columns=['Fold_{}'.format(i) for i in range(1,N+1)
                            index=df_all.columns )
fprs, tprs, scores = [], [], []

skf = StratifiedKFold(n_split=N, random_state=N, shuffle=True)
```

```
for fold, (trn_idx, val_idx) in enumerate(skf.split(X_train, y_train), 1):

    print(f'Fold {fold}')

    # Fitting the model.
    leaderboard_model.fit(X_train[trn_idx], y_train[val_idx])

    # Computing Train AUC score.
    trn_fpr, trn_tpr, trn_thresholds = roc_curve(y_train[trn_idx],
                                                 leaderboard_model.predict_proba
                                                 (X_train[trn_idx])[:, 1])
    trn_auc_score = auc(trn_fpr, trn_tpr)
    # Computing validation AUC score.
    val_fpr, val_tpr, val_thresholds = roc_curve(y_train[val_idx],
                                                 leaderboard_model.predict_proba
                                                 (X_train[val_idx])[:, 1])
    val_auc_score = auc(val_fpr, val_tpr)

    scores.append((trn_auc_score, val_auc_score))
    fprs.append(val_fpr)
    tprs.append(val_tpr)

    # X_test prob.                                  format(fold)
    probs.loc[:, 'fold_{}_prob_0'] = leaderboard_model.predict_proba(
                                                 X_test)[:, 0])

    #                    |
    importances.iloc[:, fold-1] = leaderboard_model.feature_importances_
```

```
        oob += leaderboard_model.oob_score_ / N
        print(f'Fold {fold} OOB score: {leaderboard_model.oob_score_}')

print(f'Ave OOB score: {oob}')
```

3.2 Feature Importances

```
importances['Mean_importances'] = importances.mean(axis=1)
importances.sort_values(by='Mean_importances', ascending=False,
                        inplace=True)

sns.barplot(x='Mean_importances', y=importances.index, data=importances)
plt.show()
```